

prevAll — demo script

a terminal cockpit for your life domains. v1.6.2

0 · the pitch (30s)

- personal cockpit that runs council deliberations across Claude, Codex, Antigravity, Ollama
- every domain (career, wealth, tax, health, business, ...) is a folder; every council verdict is logged
- ships with bundled 21-question canonical benchmark to grade new models against your real decisions
- single binary, two-machine sync over Tailscale, no cloud lock-in

1 · install & first boot (1 min)

- `curl -fsSL prevail.sh/install | bash` — pulls v1.6.2 tarball
- tarball contains: prevail binary + vault-demo + AGENTS-operating.md
- prevail — first run shows wizard
- wizard offers three paths:
 - use bundled demo vault (20 seeded domains + 21 bench questions)
 - point at existing vault path
 - type custom path (gets created)
- detected CLIs are auto-probed (claude, codex, agy, ollama)

2 · the cockpit (1 min)

- left rail: 20 domain chips, gold for active, dim for the rest
- top banner: `◇ configure ◇ bench > tools chips`
- main pane: chat | state | quickstart | prompts | skills | journal | calibration tabs
- tab cycles views, `↑↓` moves domains, click anywhere works
- footer hint: talk to claude · [e] edit · [r] refresh · [q] quit

3 · talk to a domain (1 min)

- select **wealth** in the rail, type a question in chat
- reply streams from claude (or codex / agy / ollama depending on default)
- every exchange logs to `wealth/_log/<date>.md` — greppable forever
- switch view to **state** — sees the running state.md, open loops, last verdict

4 · council mode (1 min)

- press `Shift+C` on a hard question to convene the council
- claude + codex + agy answer in parallel, chair model synthesizes
- verdict written as `## HH:MM · ■ council` section in the day's log
- next morning's chief brief surfaces last council verdict across all domains

5 · the benchmark — the flagship demo (5 min)

press Shift+B or click  bench on the banner

5a · question list (top of panel)

- 21 canonical questions across 8 capability categories
- each row:  <question-id> · <domain>
 - click any row → expands inline to show prompt + context + expected decision + verdict keywords + attachments
 - (new in v1.6.2 — lets you audit the benchmark without opening every .md file)
- categories: binary decisions, document analysis (5 questions with attachments), recency, cultural nuance, brevity, insufficient-info, bias detection, complex tax trap

5b · customize (under question list)

- + new question — wizard-driven scaffold
- > import from journal — pulls a recent council verdict from your real decision log
- \ edit highlighted — opens the .md in \$EDITOR
- > open folder — reveals vault/benchmark/ in Finder

5c · pick models (multi-select, cross-provider)

- tab between Claude / Codex / Antigravity / Ollama
- check multiple models per CLI (e.g. opus-4-7 + opus-4-6 + haiku)
- selections persist across tabs — check 2 from Claude + 3 from Codex
- header row shows: across: 6 total · Claude 2 + Codex 3 + Antigravity 1
- > run 6 models × 21 questions = 126 calls

5d · live run

- progress strip per active model: ✓ question-id 1.2s
- after each model finishes, its score lands and the next starts
- judge model is the same CLI as the runner (no cross-provider judging bias)

5e · per-model results (click-to-drill)

- summary row per model: keyword_avg: 46% judge_avg: 9.7 / 10
- per-question rows: > question-id 60% 10/10 rationale...
- click any row → expands to show:
 - the prompt the model received
 - the model's full reply
 - expected decision (in green)
 - keyword hits / misses
 - judge rationale (full)
- (v1.6.1 fixed wrap-to-col-0 bug so rows stay single-line at any width)

5f · leaderboard (bottom of panel, click any row)

- every past run sorted by judge_avg
- click any row → hydrates the results section above with that run's saved data
 - no re-firing of the council, no re-spending
 - browse historical runs for any model you've ever tested
 -  viewing saved run: <label> badge tells you what you're looking at
 - (new in v1.6.1)

6 · attachment-driven questions (2 min)

- 5 of the 21 questions reference long-form documents
- attachments inlined into the prompt with [ATTACHMENT: name]...[END] delimiters
 - uniform across Claude / Codex / Agy / Ollama — no per-CLI attachment syntax
 - 16,000 char/file budget keeps prompt size bounded
- the 5 attachment scenarios:
 - **wealth-q4-pnl-review** — Q4 statements with AR aging drift (31→44 days, the load-bearing signal)
 - **tax-end-of-year-strategy** — draft 1040 + W-2 + 1099-DIV + brokerage; tests harvest + mega-backdoor + IRMAA trap
 - **health-lab-panel-review** — full panel; tests prioritizing signal over abundance of yellow flags
 - **legal-msa-review** — consulting MSA with uncapped indemnification, 18-mo non-compete
 - **business-pricing-decision** — pricing tiers + competitor scan + Pinetree at 67-day AR
- bash scripts/build-bench-attachments.sh generates companion PDFs for visual review

7 · the 6 scenario probes

- **recency-knowledge-cutoff** — tests if the model knows about June 18 2026 gemini→agy transition
- **family-cultural-obligation** — Nigerian immigrant sister's wedding \$40k ask; tests cultural literacy vs Western individualist default
- **brevity-test** — one-sentence-only format compliance
- **tax-passive-loss-trap** — \$50k K-1 syndication loss vs \$280k W-2; IRC §469 trap (most LLMs say YES, wrong)
- **insufficient-info-recognition** — HDHP vs PPO; correct answer is REFUSE to recommend, ask for the missing facts
- **bias-trap-anchoring** — Patek "\$10k below retail" anchor; correct answer rejects the frame entirely

8 · expected demo result (talking point)

- Claude opus 4.7 leaderboard topper: judge 9.7 / 10 · keyword 46%
- haiku at judge 8.3 / 10 · keyword 34% — visible quality delta
- codex models score lower on cultural / brevity (use as differentiator)
- ollama local models show their tradeoff in real time without leaving the terminal

9 · the rest of the cockpit (rapid fire)

- tools panel: configure MCP, telegram gateway (OpenClaw), keys, themes
- MCP server: exposes vault to external Claude Desktop or Claude Code as a tool source
- telegram: chat with your council from any device via OpenClaw bridge
- prevail upgrade: one command, .tar.gz substring match, auto-extract, no manual steps
- ~/.ai/ shared knowledge system — Paperclip (30 agents) + OpenClaw read the same STATE.md files
- two-machine sync: Mac Mini is source of truth, MBP symlinks via Tailscale, zero sync code

10 · close (30s)

- everything lives in your filesystem as markdown — no proprietary database, no cloud account
- the benchmark is your moat: as you import real verdicts via --from-log, new models get graded against decisions you already made
- every release ships full vault-demo so new users get a working cockpit in 60 seconds
- shipped through this session: **v0.9.0** → **v1.6.2** in one sitting (15 releases, 0 regressions)